

Box Office Prediction-Copy1

October 25, 2018

0.1 Initial imports and loading data with Pandas

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import tree
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
%matplotlib inline

pd.set_option('mode.chained_assignment', None)
pd.set_option('display.float_format', '{:,.2f}'.format)
```

```
In [3]: data = pd.read_csv('movie_metadata.csv')
```

0.2 Taking a look at the data

You need to "run" the two cells below, to do that select the cell and press: *Shift-Enter*

```
In [4]: # Run this cell (to do so press Shift-Enter)
data.head(5)
```

```
Out[4]:
```

	color	director_name	num_critic_for_reviews	duration	\
0	Color	James Cameron	723.00	178.00	
1	Color	Gore Verbinski	302.00	169.00	
2	Color	Sam Mendes	602.00	148.00	
3	Color	Christopher Nolan	813.00	164.00	
4	NaN	Doug Walker	nan	nan	

	director_facebook_likes	actor_3_facebook_likes	actor_2_name	\
0	0.00	855.00	Joel David Moore	
1	563.00	1,000.00	Orlando Bloom	
2	0.00	161.00	Rory Kinnear	
3	22,000.00	23,000.00	Christian Bale	
4	131.00	nan	Rob Walker	

	actor_1_facebook_likes	gross	genres	\
0	1,000.00	760,505,847.00	Action Adventure Fantasy Sci-Fi	
1	40,000.00	309,404,152.00	Action Adventure Fantasy	
2	11,000.00	200,074,175.00	Action Adventure Thriller	
3	27,000.00	448,130,642.00	Action Thriller	
4	131.00	nan	Documentary	

	...	num_user_for_reviews	language	country	content_rating	\
0	...	3,054.00	English	USA	PG-13	
1	...	1,238.00	English	USA	PG-13	
2	...	994.00	English	UK	PG-13	
3	...	2,701.00	English	USA	PG-13	
4	...	nan	NaN	NaN	NaN	

	budget	title_year	actor_2_facebook_likes	imdb_score	aspect_ratio	\
0	237,000,000.00	2,009.00	936.00	7.90	1.78	
1	300,000,000.00	2,007.00	5,000.00	7.10	2.35	
2	245,000,000.00	2,015.00	393.00	6.80	2.35	
3	250,000,000.00	2,012.00	23,000.00	8.50	2.35	
4	nan	nan	12.00	7.10	nan	

	movie_facebook_likes
0	33000
1	0
2	85000
3	164000
4	0

[5 rows x 28 columns]

In [5]: data.shape

Out[5]: (5043, 28)

In [6]: data.describe()

Out[6]:

	num_critic_for_reviews	duration	director_facebook_likes	\
count	4,993.00	5,028.00	4,939.00	
mean	140.19	107.20	686.51	
std	121.60	25.20	2,813.33	
min	1.00	7.00	0.00	
25%	50.00	93.00	7.00	
50%	110.00	103.00	49.00	
75%	195.00	118.00	194.50	
max	813.00	511.00	23,000.00	

	actor_3_facebook_likes	actor_1_facebook_likes	gross	\
count	5,020.00	5,036.00	4,159.00	

mean	645.01	6,560.05	48,507,385.63
std	1,665.04	15,020.76	68,471,915.43
min	0.00	0.00	162.00
25%	133.00	614.00	5,351,178.00
50%	371.50	988.00	25,528,495.00
75%	636.00	11,000.00	62,319,957.00
max	23,000.00	640,000.00	760,505,847.00

	num_voted_users	cast_total_facebook_likes	facenumber_in_poster \
count	5,043.00	5,043.00	5,030.00
mean	83,668.16	9,699.06	1.37
std	138,485.26	18,163.80	2.01
min	5.00	0.00	0.00
25%	8,593.50	1,411.00	0.00
50%	34,359.00	3,090.00	1.00
75%	96,309.00	13,756.50	2.00
max	1,689,764.00	656,730.00	43.00

	num_user_for_reviews	budget	title_year \
count	5,022.00	4,551.00	4,935.00
mean	272.77	39,752,620.44	2,002.47
std	377.98	206,114,898.45	12.47
min	1.00	218.00	1,916.00
25%	65.00	6,000,000.00	1,999.00
50%	156.00	20,000,000.00	2,005.00
75%	326.00	45,000,000.00	2,011.00
max	5,060.00	12,215,500,000.00	2,016.00

	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes
count	5,030.00	5,043.00	4,714.00	5,043.00
mean	1,651.75	6.44	2.22	7,525.96
std	4,042.44	1.13	1.39	19,320.45
min	0.00	1.60	1.18	0.00
25%	281.00	5.80	1.85	0.00
50%	595.00	6.60	2.35	166.00
75%	918.00	7.20	2.35	3,000.00
max	137,000.00	9.50	16.00	349,000.00

Some key points from this table: - Avg movie duration is 107.2 minutes - Avg imdb is 6.44 - Avg number of users reviews is 272

0.3 Cleaning the data

0.3.1 Dealing with duplicates

```
In [10]: print ('Number of duplicates in data: {}'.format(
            sum(data.duplicated(subset=['movie_title', 'title_year'], keep=False))))
```

Number of duplicates in data: 0

```
In [9]: data = data.drop_duplicates(subset=['movie_title', 'title_year'], keep='first').copy()
```

0.3.2 Fixing Null and some zero values

```
In [11]: # check if data has any null/nan values
         data.isnull().values.any()
```

```
Out[11]: True
```

```
In [15]: # Check how many values are null in each column
         def show_missing_data(data):
             missing_data = data.isnull().sum().reset_index()
             missing_data.columns = ['column_name', 'missing_count']
             missing_data['filling_factor'] = (data.shape[0] - missing_data['missing_count']) /
             return missing_data.sort_values('filling_factor').reset_index(drop=True)

         show_missing_data(data)[:5]
```

```
Out[15]:
```

	column_name	missing_count	filling_factor
0	budget	266	93.44
1	aspect_ratio	104	97.44
2	content_rating	64	98.42
3	plot_keywords	40	99.01
4	actor_3_name	13	99.68

As we are working with the Gross Box Office, rows without it are of no use. So we will exclude those films that are missing the Gross Box Office.

```
In [16]: data.dropna(subset=['gross'], how='all', inplace=True)
         show_missing_data(data)[:5]
```

```
Out[16]:
```

	column_name	missing_count	filling_factor
0	budget	266	93.44
1	aspect_ratio	104	97.44
2	content_rating	64	98.42
3	plot_keywords	40	99.01
4	actor_3_name	13	99.68

Fill out missing budget datapoints with the median budget for the year it was released.

```
In [386]: median_budget_per_year = data.groupby('title_year')['budget'].transform('median')
         data['budget'].fillna(median_budget_per_year, inplace=True)

         show_missing_data(data)[:5]
```

```
Out[386]:
```

	column_name	missing_count	filling_factor
0	aspect_ratio	104	97.44
1	content_rating	64	98.42
2	plot_keywords	40	99.01
3	actor_3_name	13	99.68
4	actor_3_facebook_likes	13	99.68

Fill out the rest of the missing data

```
In [387]: data.fillna(0, inplace=True)
```

Delete all rows where title_year is zero

```
In [388]: data = data[data['title_year'] != 0]
```

Budgets are in each country's currency so we are going to use only US movies


```
In [389]: data = data[data['country'] == 'USA']
```

0.4 Understanding the data

```
In [25]: # IMDb rating distribution
data.hist(column='imdb_score')
```

```
Out[25]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fa49f313c88>]], dtype=object)


/opt/conda/lib/python3.5/site-packages/matplotlib/font_manager.py:1297: UserWarning: findfont: F
(prop.get_family(), self.defaultFamily[fonttext]))
```



```
In [391]: # Movies per year
```

```
data.hist(column='title_year')
```

```
Out[391]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x129be56d8>]], dtype=object)
```



```
In [31]: # Median gross box office per actor
```

```
fig = plt.figure(figsize=(8,8))
```

```
comparison_df = data.groupby('actor_1_name', as_index=False).mean().sort_values('gross')
```

```
name_count_key = data['actor_1_name'].value_counts().to_dict()
```

```
comparison_df['films'] = comparison_df['actor_1_name'].map(name_count_key)
```

```
comparison_df['actor_1_name'] = comparison_df['actor_1_name'].map(str) + " (" + comparison_df['films']
```

```
comparison_df[comparison_df['films'] >= 5][['actor_1_name', 'gross']][10::-1].set_index('gross')
```

```
plt.legend().set_visible(False)
```


```
plt.title("Median Gross of Actor_1_name's Films")
```

```
plt.ylabel("Actor_1_name (# films)")
```

```
plt.xlabel("Gross (in million)")
```


```
Out[31]: <matplotlib.text.Text at 0x7fa49ef14eb8>
```

```
/opt/conda/lib/python3.5/site-packages/matplotlib/font_manager.py:1297: UserWarning: findfont: F
(prop.get_family(), self.defaultFamily[fonttext]))
```



```
In [393]: # title year vs gross
data.plot.scatter(x='title_year', y='gross')
```


```
Out[393]: <matplotlib.axes._subplots.AxesSubplot at 0x12aa7a9b0>
```



In [394]: *# budget vs gross*

```
data.plot.scatter(x='budget', y='gross')
```

Out[394]: <matplotlib.axes._subplots.AxesSubplot at 0x12b46a9e8>




```
In [32]: data.corr()
```

```
Out[32]:
```

	num_critic_for_reviews	duration \
num_critic_for_reviews	1.00	0.28
duration	0.28	1.00
director_facebook_likes	0.19	0.21
actor_3_facebook_likes	0.28	0.14
actor_1_facebook_likes	0.18	0.10
gross	0.49	0.29
num_voted_users	0.61	0.37
cast_total_facebook_likes	0.25	0.14
facenumber_in_poster	-0.03	0.01
num_user_for_reviews	0.58	0.37
budget	0.49	0.30
title_year	0.39	-0.11
actor_2_facebook_likes	0.28	0.15
imdb_score	0.36	0.38
aspect_ratio	0.24	0.18
movie_facebook_likes	0.70	0.25

	director_facebook_likes	actor_3_facebook_likes \
num_critic_for_reviews	0.19	0.28
duration	0.21	0.14
director_facebook_likes	1.00	0.13
actor_3_facebook_likes	0.13	1.00
actor_1_facebook_likes	0.09	0.25
gross	0.14	0.28
num_voted_users	0.32	0.27
cast_total_facebook_likes	0.12	0.47
facenumber_in_poster	-0.05	0.10
num_user_for_reviews	0.25	0.22
budget	0.10	0.27
title_year	-0.06	0.12
actor_2_facebook_likes	0.12	0.54
imdb_score	0.22	0.09
aspect_ratio	0.06	0.07
movie_facebook_likes	0.18	0.30

	actor_1_facebook_likes	gross	num_voted_users \
num_critic_for_reviews	0.18	0.49	0.61
duration	0.10	0.29	0.37
director_facebook_likes	0.09	0.14	0.32
actor_3_facebook_likes	0.25	0.28	0.27
actor_1_facebook_likes	1.00	0.13	0.18
gross	0.13	1.00	0.64

num_voted_users	0.18	0.64	1.00
cast_total_facebook_likes	0.95	0.22	0.25
facenumber_in_poster	0.06	-0.03	-0.04
num_user_for_reviews	0.13	0.56	0.79
budget	0.15	0.65	0.42
title_year	0.09	0.03	0.02
actor_2_facebook_likes	0.38	0.24	0.25
imdb_score	0.12	0.26	0.50
aspect_ratio	0.07	0.13	0.14
movie_facebook_likes	0.13	0.38	0.53

	cast_total_facebook_likes	facenumber_in_poster	\
num_critic_for_reviews	0.25	-0.03	
duration	0.14	0.01	
director_facebook_likes	0.12	-0.05	
actor_3_facebook_likes	0.47	0.10	
actor_1_facebook_likes	0.95	0.06	
gross	0.22	-0.03	
num_voted_users	0.25	-0.04	
cast_total_facebook_likes	1.00	0.08	
facenumber_in_poster	0.08	1.00	
num_user_for_reviews	0.19	-0.08	
budget	0.23	-0.03	
title_year	0.12	0.08	
actor_2_facebook_likes	0.62	0.07	
imdb_score	0.13	-0.08	
aspect_ratio	0.09	0.01	
movie_facebook_likes	0.21	0.01	

	num_user_for_reviews	budget	title_year	\
num_critic_for_reviews	0.58	0.49	0.39	
duration	0.37	0.30	-0.11	
director_facebook_likes	0.25	0.10	-0.06	
actor_3_facebook_likes	0.22	0.27	0.12	
actor_1_facebook_likes	0.13	0.15	0.09	
gross	0.56	0.65	0.03	
num_voted_users	0.79	0.42	0.02	
cast_total_facebook_likes	0.19	0.23	0.12	
facenumber_in_poster	-0.08	-0.03	0.08	
num_user_for_reviews	1.00	0.42	0.02	
budget	0.42	1.00	0.22	
title_year	0.02	0.22	1.00	
actor_2_facebook_likes	0.20	0.24	0.12	
imdb_score	0.34	0.07	-0.14	
aspect_ratio	0.15	0.20	0.12	
movie_facebook_likes	0.40	0.33	0.28	

	actor_2_facebook_likes	imdb_score	aspect_ratio	\
--	------------------------	------------	--------------	---

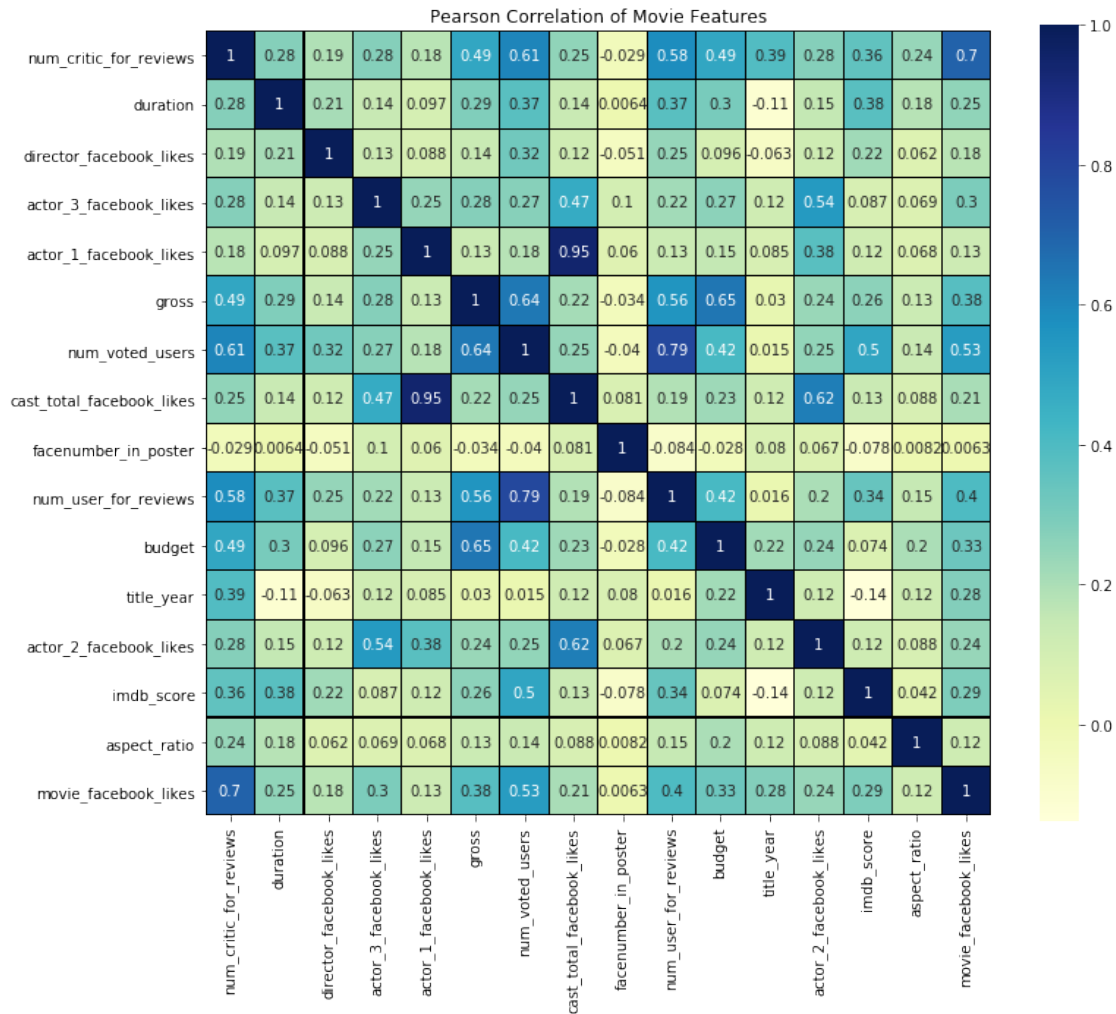
num_critic_for_reviews	0.28	0.36	0.24
duration	0.15	0.38	0.18
director_facebook_likes	0.12	0.22	0.06
actor_3_facebook_likes	0.54	0.09	0.07
actor_1_facebook_likes	0.38	0.12	0.07
gross	0.24	0.26	0.13
num_voted_users	0.25	0.50	0.14
cast_total_facebook_likes	0.62	0.13	0.09
facenumber_in_poster	0.07	-0.08	0.01
num_user_for_reviews	0.20	0.34	0.15
budget	0.24	0.07	0.20
title_year	0.12	-0.14	0.12
actor_2_facebook_likes	1.00	0.12	0.09
imdb_score	0.12	1.00	0.04
aspect_ratio	0.09	0.04	1.00
movie_facebook_likes	0.24	0.29	0.12

	movie_facebook_likes
num_critic_for_reviews	0.70
duration	0.25
director_facebook_likes	0.18
actor_3_facebook_likes	0.30
actor_1_facebook_likes	0.13
gross	0.38
num_voted_users	0.53
cast_total_facebook_likes	0.21
facenumber_in_poster	0.01
num_user_for_reviews	0.40
budget	0.33
title_year	0.28
actor_2_facebook_likes	0.24
imdb_score	0.29
aspect_ratio	0.12
movie_facebook_likes	1.00

```
In [396]: # Set up the matplotlib figure
          f, ax = plt.subplots(figsize=(12, 10))
          plt.title('Pearson Correlation of Movie Features')

          # Draw the heatmap using seaborn
          sns.heatmap(data.corr(),linewidths=0.25,vmax=1.0, square=True, cmap="YlGnBu", linecolor=

Out[396]: <matplotlib.axes._subplots.AxesSubplot at 0x130f627f0>
```



As we can see from the heatmap, there are regions (features) where we can see quite positive linear correlations amongst each other, given the darker shade of the colours - top left-hand corner and bottom right quarter. This is a good sign as it means we may be able to find linearly correlated features for which we can perform PCA projections on.

```
In [33]: data.corr()['gross'].sort_values(ascending=False)
```

```
Out[33]: gross          1.00
         budget         0.65
         num_voted_users 0.64
         num_user_for_reviews 0.56
         num_critic_for_reviews 0.49
         movie_facebook_likes 0.38
         duration        0.29
         actor_3_facebook_likes 0.28
         imdb_score       0.26
         actor_2_facebook_likes 0.24
```

```

cast_total_facebook_likes    0.22
director_facebook_likes      0.14
actor_1_facebook_likes       0.13
aspect_ratio                 0.13
title_year                   0.03
facenumber_in_poster        -0.03
Name: gross, dtype: float64

```

The gross box office correlates strongly with num_voted_users, num_users_for_reviews and movie_facebook_likes. But some of those features are also highly correlated among each other (as you can see in the heatmap above).

0.5 Gross Box Office Prediction

0.5.1 Getting numerical data

```

In [34]: numerical_columns = data.dtypes[data.dtypes != 'object'].index
         numerical_data = data[numerical_columns]

```

```

# we drop aspect_ratio, as it doesn't provide any useful info
numerical_data.drop('aspect_ratio', axis=1, inplace=True)
numerical_data.head(3)

```

```

Out[34]:
   num_critic_for_reviews  duration  director_facebook_likes  \
0                723.00    178.00                0.00
1                302.00    169.00               563.00
3                813.00    164.00            22,000.00

   actor_3_facebook_likes  actor_1_facebook_likes    gross  \
0                855.00            1,000.00  760,505,847.00
1               1,000.00            40,000.00  309,404,152.00
3              23,000.00            27,000.00  448,130,642.00

   num_voted_users  cast_total_facebook_likes  facenumber_in_poster  \
0            886204                4834                0.00
1            471220                48350                0.00
3           1144337               106759                0.00

   num_user_for_reviews    budget  title_year  actor_2_facebook_likes  \
0            3,054.00  237,000,000.00    2,009.00                936.00
1            1,238.00  300,000,000.00    2,007.00               5,000.00
3            2,701.00  250,000,000.00    2,012.00            23,000.00

   imdb_score  movie_facebook_likes
0          7.90                33000
1          7.10                 0
3          8.50               164000

```

0.5.2 Preparing train and test datasets

```
In [35]: train, test = train_test_split(numerical_data, test_size=0.2)
        target_train = train.pop('gross')
        target_test = test.pop('gross')
```

```
In [36]: print('Train data: {} / {} = {}'.format(len(train), len(numerical_data), float(len(train)/len(numerical_data))))
        print('Test data: {} / {} = {}'.format(len(test), len(numerical_data), float(len(test)/len(numerical_data))))
```

Train data: 2523 / 3154 = 0.7999365884590995

Test data: 631 / 3154 = 0.20006341154090043

0.5.3 Linear Regression

```
In [37]: model = LinearRegression()
        model.fit(train, target_train)
```

```
Out[37]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [38]: prediction = model.predict(test)
```

```
In [411]: # The mean squared error
        print("Mean squared error: %.2f" % mean_squared_error(target_test, prediction))


        # Explained variance score: 1 is perfect prediction
        print('Variance score: %.2f' % r2_score(target_test, prediction))

        # Plot outputs
        plt.scatter(target_test, prediction, color='black')
        # plt.plot(test, prediction, color='blue', linewidth=3)

        plt.show()
```

Mean squared error: 1922173864183296.50

Variance score: 0.62



0.5.4 Random Forest

```
In [47]: forest = RandomForestRegressor(
        max_depth=25,
        min_samples_split=15,
        n_estimators=1000,
        random_state=1)
```

```
forest.fit(train, target_train)
```

```
Out[47]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=25,
        max_features='auto', max_leaf_nodes=None,
        min_impurity_split=1e-07, min_samples_leaf=1,
        min_samples_split=15, min_weight_fraction_leaf=0.0,
        n_estimators=1000, n_jobs=1, oob_score=False, random_state=1,
        verbose=0, warm_start=False)
```

```
In [48]: forest.feature_importances_
```

```
Out[48]: array([ 0.09460968,  0.0574659 ,  0.05838327,  0.03151676,  0.04906885,
        0.00854848,  0.58500998,  0.06361421,  0.05178287])
```

```
In [49]: forest_prediction = forest.predict(test)
```

```
In [65]: # The mean squared error
print("Mean squared error: %.2f" % mean_squared_error(target_test, forest_prediction))

# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % r2_score(target_test, forest_prediction))

# Plot outputs
plt.scatter(target_test, forest_prediction, color='black')
# plt.plot(test, prediction, color='blue', linewidth=3)


plt.show()
```

Mean squared error: 3009956585804540.50

Variance score: 0.48

/opt/conda/lib/python3.5/site-packages/matplotlib/font_manager.py:1297: UserWarning: findfont: F

```
(prop.get_family(), self.defaultFamily[fonttext]))
```



0.6 Dropping post-fact data

There are post-fact variables in our data set making the prediction more accurate. Things like `num_voted_users` and `num_user_for_reviews` are after the fact metrics, so probably not as useful for prediction.


```
In [57]: train.head(2)
```

```
Out[57]:
```

	duration	director_facebook_likes	actor_3_facebook_likes	\
4462	95.00	12,000.00	636.00	
3467	95.00	213.00	92.00	

	actor_1_facebook_likes	cast_total_facebook_likes	facenumber_in_poster	\
4462	12,000.00	14420	1.00	
3467	225.00	791	0.00	

	budget	title_year	actor_2_facebook_likes
4462	1,300,000.00	1,996.00	680.00
3467	4,000,000.00	1,983.00	174.00

```
In [66]: train.drop(['num_critic_for_reviews', 'num_voted_users', 'num_user_for_reviews', 'imdb_score', 'movie_facebook_likes'], axis=1)
test.drop(['num_critic_for_reviews', 'num_voted_users', 'num_user_for_reviews', 'imdb_score', 'movie_facebook_likes'], axis=1)
train.head(2)
```

```
-----
ValueError                                Traceback (most recent call last)
```

```
<ipython-input-66-4f4aae4c6783> in <module>()
----> 1 train.drop(['num_critic_for_reviews', 'num_voted_users', 'num_user_for_reviews', 'imdb_score', 'movie_facebook_likes'], axis=1)
      2 test.drop(['num_critic_for_reviews', 'num_voted_users', 'num_user_for_reviews', 'imdb_score', 'movie_facebook_likes'], axis=1)
      3 train.head(2)

/opt/conda/lib/python3.5/site-packages/pandas/core/generic.py in drop(self, labels, axis, level, errors)
1905         new_axis = axis.drop(labels, level=level, errors=errors)
1906     else:
-> 1907         new_axis = axis.drop(labels, errors=errors)
1908         dropped = self.reindex(**{axis_name: new_axis})
1909     try:

/opt/conda/lib/python3.5/site-packages/pandas/indexes/base.py in drop(self, labels, errors)
3260         if errors != 'ignore':
3261             raise ValueError('labels %s not contained in axis' %
-> 3262                               labels[mask])
3263         indexer = indexer[~mask]
3264         return self.delete(indexer)
```

```
ValueError: labels ['num_critic_for_reviews' 'num_voted_users' 'num_user_for_reviews' 'imdb_score' 'movie_facebook_likes'] not contained in axis
```

```

In [41]: pre_data_forest = RandomForestRegressor(
        max_depth=25,
        min_samples_split=15,
        n_estimators=1000,
        random_state=1)

pre_data_forest.fit(train, target_train)

Out[41]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=25,
        max_features='auto', max_leaf_nodes=None,
        min_impurity_split=1e-07, min_samples_leaf=1,
        min_samples_split=15, min_weight_fraction_leaf=0.0,
        n_estimators=1000, n_jobs=1, oob_score=False, random_state=1,
        verbose=0, warm_start=False)

In [415]: second_prediction = pre_data_forest.predict(test)

In [416]: # The mean squared error
        print("Mean squared error: %.2f" % mean_squared_error(target_test, second_prediction))


        # Explained variance score: 1 is perfect prediction
        print('Variance score: %.2f' % r2_score(target_test, second_prediction))

        # Plot outputs
        plt.scatter(target_test, second_prediction, color='black')
        # plt.plot(test, prediction, color='blue', linewidth=3)

        plt.show()

Mean squared error: 3013829195927549.50
Variance score: 0.41

```



0.7 Over/Under performing movies

```
In [417]: numerical_data_target = numerical_data.pop('gross')
          all_data_prediction = forest.predict(numerical_data)
```

```
In [484]: performance_df = data.copy()
```

```
performance_df["prediction"] = all_data_prediction
performance_df["performance_diff"] = numerical_data_target - all_data_prediction
```

```
performance_df.sort_values(['performance_diff'], ascending=False, inplace=True)
```

```
In [485]: ind = np.arange(5)
          width = 0.35
```

```
fig, ax = plt.subplots(figsize=(12, 8))
```


```
gross = ax.bar(ind, performance_df.gross[:5], width, color='r')
predicted_gross = ax.bar(ind + width, performance_df.prediction[:5], width, color='g')
```

```
plt.title("Overperforming movies")
plt.ylabel("Gross (in million)")
```

```
ax.set_xticks(ind + width / 2)
```

```
ax.set_xticklabels(performance_df.movie_title[:5], rotation='vertical')
ax.legend((gross[0], predicted_gross[0]), ('Gross', 'Prediction'))
```

Out[485]: <matplotlib.legend.Legend at 0x130e9be10>



```
In [59]: performance_repr[:-6:-1]
```

```
ind = np.arange(5)
width = 0.35
```

```
fig, ax = plt.subplots(figsize=(12, 8))
```

```
gross = ax.bar(ind, performance_df.gross[:-6:-1], width, color='r')
```

```
predicted_gross = ax.bar(ind + width, performance_df.prediction[:-6:-1], width, color='g')
```

```
plt.title("Overperforming movies")
```

```
plt.ylabel("Gross (in million)")

ax.set_xticks(ind + width / 2)
ax.set_xticklabels(performance_df.movie_title[:-6:-1], rotation='vertical')
ax.legend((gross[0], predicted_gross[0]), ('Gross', 'Prediction'))
```

NameError Traceback (most recent call last)

```
<ipython-input-59-42fe547378f6> in <module>()
----> 1 performance_repr[:-6:-1]
      2
      3 ind = np.arange(5)
      4 width = 0.35
      5
```

NameError: name 'performance_repr' is not defined

In []: